


# CPU Scheduling

Scheduling criteria, Scheduling Algorithms: Preemptive and non-preemptive, FCFS, SJF, RR

# CPU-I/O Burst Cycle

- Process execution consists of a cycle of **CPU execution** and **I/O wait**.
- When one process does I/O, a scheduler will typically switch the CPU to another process.



# CPU-Bound and I/O-Bound Processes

- A typical process execution has
  - a large number of short CPU bursts
  - a small number of long CPU bursts
- A **CPU-bound** process – mostly long CPU bursts
- A **I/O-bound** process – mostly short CPU bursts

# CPU Scheduling

- Selecting a waiting process from the ready queue and allocating the CPU to it.
  - It executes frequently
  - It must be fast
    - If it executes every 100 ms and it takes 10 ms, then 10% of the CPU time has been used for scheduling.
- Many ways to do CPU scheduling.  
How do we compare them?


# Scheduling Criteria

- Maximize
  - **CPU utilization** – percentage of CPU time spent doing work
  - **Throughput** – number of processes completed per time unit.
- Minimize
  - **Turnaround time** – interval from submission to completion of a process.
  - **Waiting time** – time spent waiting in the ready queue
  - **Response time** – time from submission of a request until the first response is produced (relevant for interactive systems)

# Preemptive vs. Non-preemptive Scheduling

- **Non-preemptive Scheduling:** Once the system has assigned a CPU to a process, the system cannot remove that CPU from that process
  - Simpler
  - Up to the process to release the CPU
- **Preemptive Scheduling:** The system can remove the CPU from the running process.
  - Need extra hardware (timer)
  - What if the process is in the middle of updating some data?

# Scheduling Algorithms



Which process should I pick?



CPU Scheduler

# First-Come, First-Served (FCFS)

- Use **FIFO** queue.
  - FIFO = First-In First-Out
- **Non-preemptive** – A process doesn't give up CPU until it either terminates or performs I/O.



# FCFS Example

- The following set of processes arrive at time 0

Process	CPU Burst Time (ms)
P1	24
P2	3
P3	3

Assume the processes arrive in the order P1, P2, P3, and are served in FCFS order

# FCFS Example

Process	CPU Burst Time
P1	24
P2	3
P3	3

Gantt Chart



Average Waiting time  $= (0\text{ms} + 24\text{ms} + 27\text{ms}) / 3 = 17\text{ms}$

Average Turnaround time  $= (24\text{ms} + 27\text{ms} + 30\text{ms}) / 3 = 27\text{ms}$

**Waiting** – time spent waiting in the ready queue

**Turnaround** – interval from submission to completion of a process.

What if the processes arrive in the order P2, P3, P1?

Process	CPU Burst
P2	3
P3	3
P1	24

Gantt Chart



Average Waiting time  $= (0\text{ms} + 3\text{ms} + 6\text{ms}) / 3 = 3\text{ms}$

Average Turnaround time  $= (30\text{ms} + 3\text{ms} + 6\text{ms}) / 3 = 13\text{ms}$

Arrival order	Average waiting time	Turnaround time
P1, P2, P3	17	27
P2, P3, P1	3	13

# Performance of FCFS Scheduling

- The average waiting time may vary if the process CPU-burst times vary greatly
  - **Convoy effect**—all the other processes wait for the one big process to release the CPU.
  - Low CPU and device utilization

# Shortest-Job-First(SJF)

- Assign CPU to the process that has the smallest next CPU-burst time
- May be either preemptive or non-preemptive

# SJF Example

- Consider the following set of processes and their arrival times

Process	Arrival Time	CPU Burst (ms)
P1	0	8
P2	0	5
P3	2	7
P4	2	2

The arrival order of the processes is P1, P2, P3, P4


# Use non-preemptive SJF



Average Waiting time =  $(14+0+5+7)/4 = 6.5\text{ms}$

Process	Arrival Time	CPU Burst (ms)
P1	0	8
P2	0	5
P3	2	7
P4	2	2


# Use FCFS



Average Waiting time =  $(0+8+13+20)/4 = 10.25\text{ms}$

Use preemptive SJF

Process	Arrival Time	CPU burst
P1	0	8
P2	0	5
P3	2	7
P4	2	2



Average Waiting time =  $(14+2+5+7)/4 = 7\text{ms}$

P1  $14-0$  (0 is arrival time) = 14

P2  $7-2$  = 5

P3  $7$  = 7

P4  $2$  = 2



# Shortest-Job-First(SJF)

- Optimal with respect to average waiting time
- How does scheduler know the length of the next CPU-burst time?

# Round-Robin(RR) Scheduling

- Similar to FCFS but with preemption.
- Have a **time quantum** (time slice). Let the first process in the queue run until it exceeds the time quantum, then run the next process.

# RR - Example

- Consider the following set of processes with their arrival times


Process	Arrival Time	CPU Burst (ms)
P1	0	24
P2	0	3
P3	0	9

The arrival order of the processes is P1, P2, P3

# Use RR Scheduling

Time quantum = 5ms

Process	Arrival Time	CPU Burst (ms)
P1	0	24
P2	0	3
P3	0	9



$$\text{Average Waiting Time} = (12+5+13)/3 = 10\text{ms}$$

Time quantum = 24ms




$$\text{Average Waiting Time} = (0+24+27)/3 = 17\text{ms}$$

Turns into FCFS scheduling

## Using RR Scheduling

Process	Arrival Time	CPU Burst (ms)
P1	0	24
P2	0	3
P3	0	9

Time quantum = 2ms



$$\text{Average waiting Time} = (12+6+13)/3 = 10.33\text{ms}$$

If context-switch time = 1ms

$$\text{Total context-switch time} = 12\text{ms}$$

# Round-Robin (RR) Scheduling

- Very small quantum
  - large context switch overhead.
- Very big quantum
  - turns into FCFS.

# Round-Robin (RR) Scheduling

- The time quantum should be large with respect to the context-switch time.
  - In most modern OSs
    - Quantum time range: 10-100 milliseconds
    - Context-switch time: <10 microseconds
- The time quantum should not be too large.
  - If the quantum time is too large, RR scheduling degenerates into FCFS.

# Summary

- CPU Scheduling
- Criteria
  - MAX: Utilization, Throughput
  - MIN: Waiting time, Turnaround time, Response time
- Algorithms
  - First-Come First-Served
  - Shortest-Job-First
  - Round-Robin