

Process Synchronization

Why Process Synchronization

Critical Section problems and solutions

Semaphores and their problems

Motivation

- Data sharing
 - Cooperating processes -- through files or messaging.
 - Threads -- directly share a logical address space
 - Concurrent access to shared data may result in data inconsistency
- Mechanisms required to ensure the orderly execution of cooperating processes or threads in order to maintain data consistency

Critical Section

- **Critical section** is a segment of code, in which there are shared modifiable data that can be accessed and modified by multiple processes (or threads).
- Critical sections can be entered by processes only in non-overlapping intervals (**mutually exclusive** in time)
 - When one process is in its critical section, other process should not be in their critical sections.

The Critical-Section Problem

- Problem: Regulate access to the critical section.
- Any solution must satisfy the following conditions:
 - **Mutual exclusion**
 - **Progress**
 - **Bounded waiting**

Conditions for Critical-Section Solutions

- **Mutual exclusion**

- If one process is executing in its critical section, then no other processes can be executing in their critical section

- **Progress**

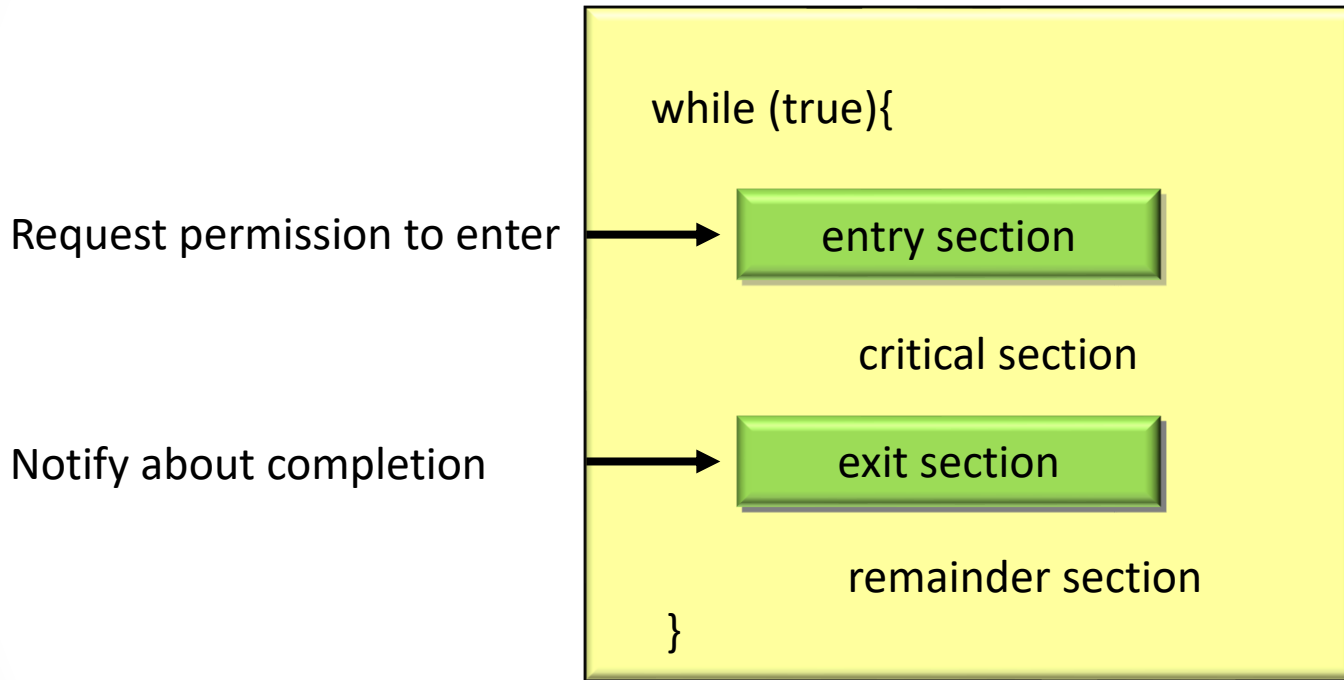
- If no process is executing in its critical section and some processes wish to enter their critical section, then one of the processes will be selected to enter its critical section.
- This selection cannot be postponed indefinitely

Conditions for Critical-Section Solutions

- **Bounded waiting:**
 - There is a **bound** on the number of times that other processes are allowed to enter their critical sections while a process is **waiting** (i.e. it has made a request to enter its critical section and that request has not been granted yet)
 - Purpose: Prevent starvation
 - **Starvation** – A process (thread) is perpetually denied access to some resources it requests.

Solving Critical-Section Problems

Generic solution



Structure of a typical process that requires access to its critical section.

Semaphores

- A **semaphore** is a mechanism provided by the system to implement **mutual exclusion**.
- A **semaphore** is
 1. a special integer variable **S** that,
 2. apart from initialization, is accessed only through two standard **atomic** operations: **acquire()** and **release()**
 3. It is associated with a queue that stores the references to the processes that are waiting.

acquire() Operation

- **acquire()** (originally termed **P()**). It is called when a process wants to enter its critical section

```
acquire(S){  
    S--;  
    if (S<0) {  
        put this process into the waiting queue  
        and block the process  
    }  
}
```

- It must be executed **indivisibly / atomically**.

release() Operation

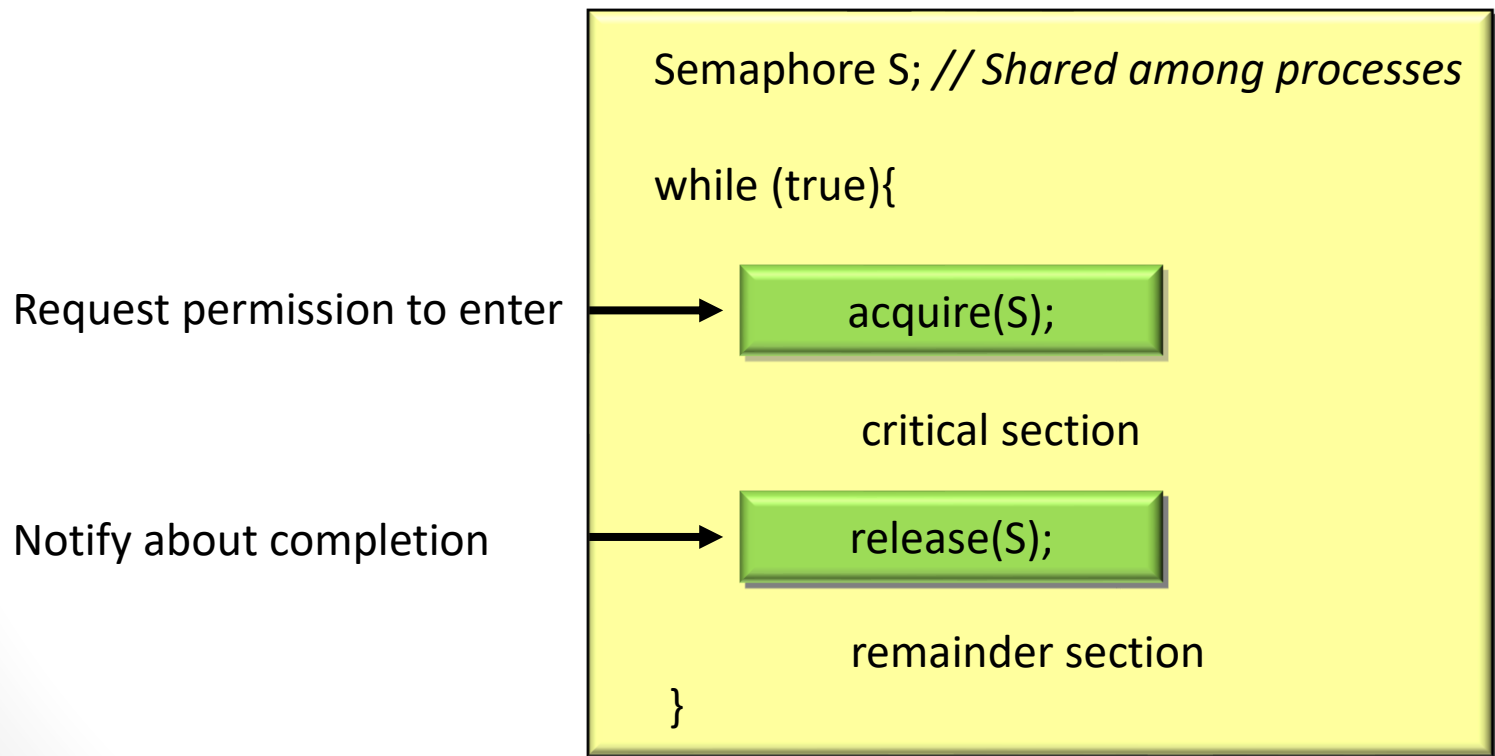
- **release()** (originally termed **V()**). It is called when a process exits its critical section

```
release(S){  
    S++;  
    if (S<=0) {  
        remove a process from the waiting queue and  
        wake up the process  
    }  
}
```

- It must be executed **indivisibly / atomically**.

Using a semaphore for mutual exclusion

- Mutual exclusion on a semaphore is enforced with **acquire()** and **release()**.



Mutex Lock and Counting Semaphores

- A **mutex lock semaphore** is used to control access to the critical section for a process
 - Its value is ..., -3, -2, -1, 0, 1. **At most 1.**
- A **counting semaphore** is used to control access to a given resource consisting a finite number of instances
 - Its value can range over an unrestricted domain

Initializing a Semaphore

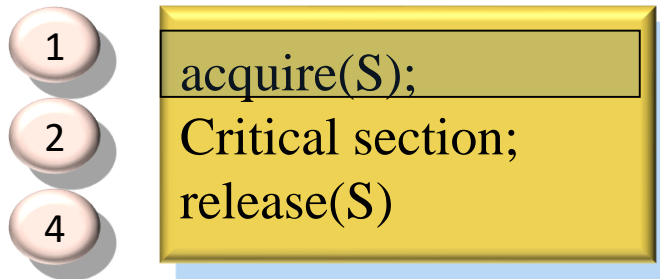
- A **mutex lock semaphore** is initialized to **1**.
- A **counting semaphore** is initialized to the number of available instances of a given resource

An Example

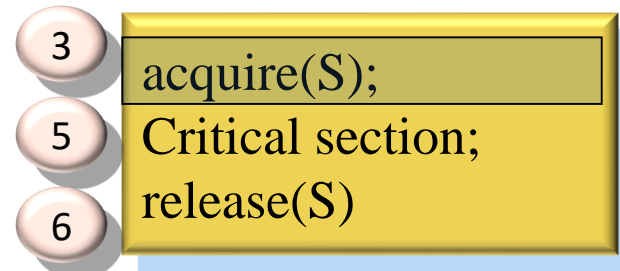
Semaphore S=1; // Shared variable initialization

S = 1

process P1



process P2



```
acquire(S){  
    S--;  
    if (S<0) {  
        put this process into queue and  
        block the process  
    }  
}
```


```
release(S){  
    S++;  
    if (S<=0) {  
        remove a process from the waiting queue  
        and wakeup the process  
    }  
}
```

Problems with Semaphores

- Use of semaphores may result in deadlock and starvation situations
 - **Deadlock**—One or more processes are waiting indefinitely for an event that can only be caused by one of the waiting processes
 - **Starvation**—Starvation occurs when a process is perpetually denied resource access. Without accessing the resources, the process is unable to complete its task.


A Deadlock Example

- Consider a system consisting of two processes P1 and P2, each accessing two semaphores, S and Q, set to the value 1:



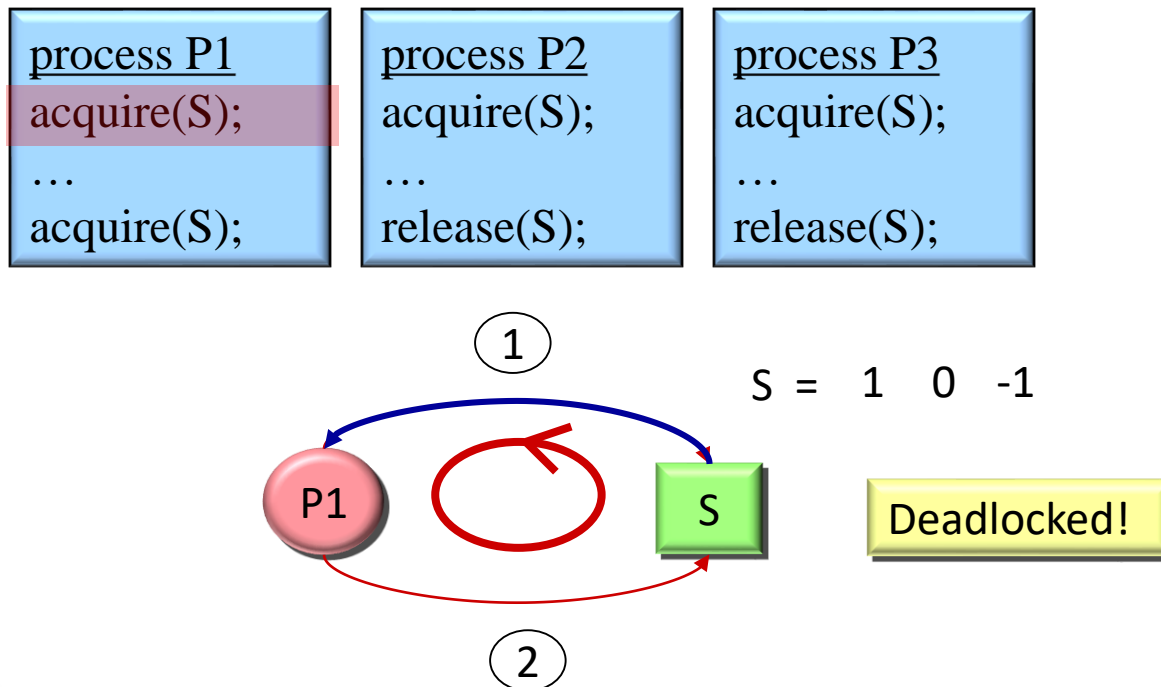
S = 1 0 -1

Q = 1 0 -1



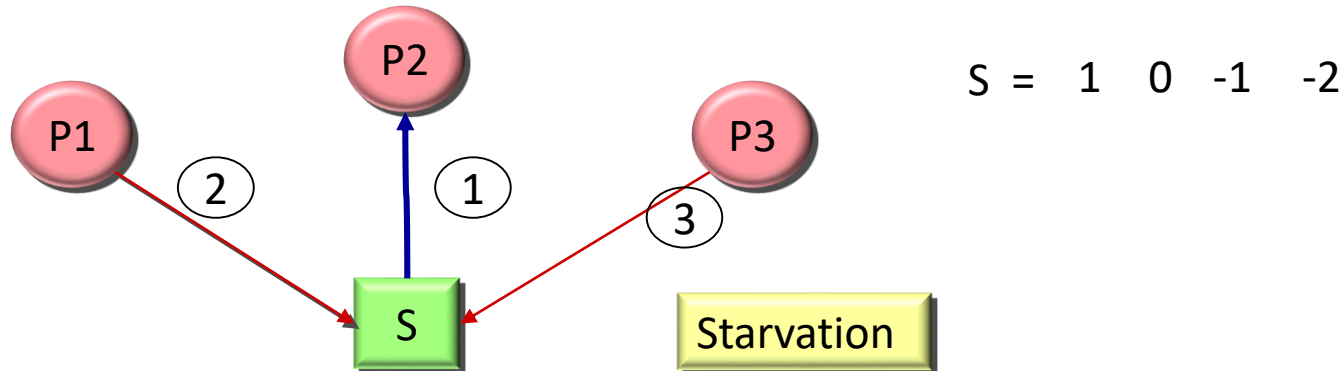
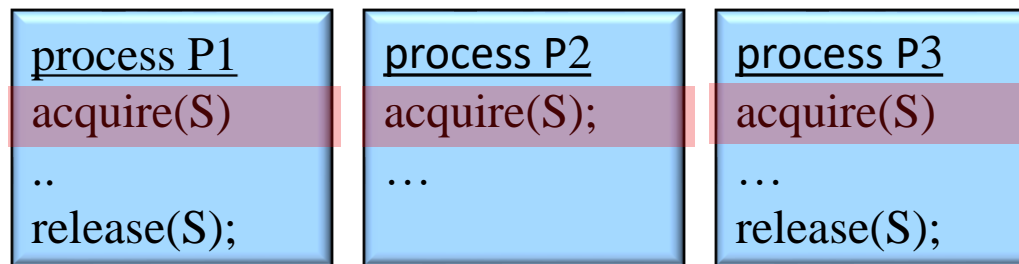
Another Deadlock Example

- Consider a system consisting of three processes P1, P2 and P3, each accessing a semaphore S which is set to the value 1




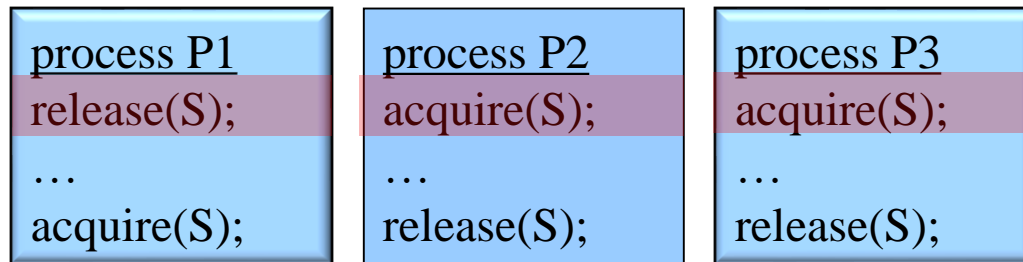
A Starvation Example

- Consider a system consisting of three processes P1, P2 and P3, each accessing a semaphore S which is set to the value 1



Examples of No Mutual Exclusion

- What happens if a process interchanges the order of acquire and release operations on a semaphore S that is set to 1?



$S = 1 \ 0 \ 1 \ 0$

Summary

- Motivation
 - Bounded-Buffer Producer and Consumer Problem
- Critical-Section Problem
 - Solution must meet three conditions
 - Mutual exclusion
 - Progress
 - Bounded waiting
- Semaphores
 - Definition
 - Implementation: `acquire()` and `release()`
 - Problems with semaphores